



Evaluating CHERI for Hardware-Assisted Memory Safety in Databases

Student:
Advisor:
Supervisor:

Cristian Sandu
Martin Fink & Ilya Meignan--Masson
Prof. Pramod Bhatotia

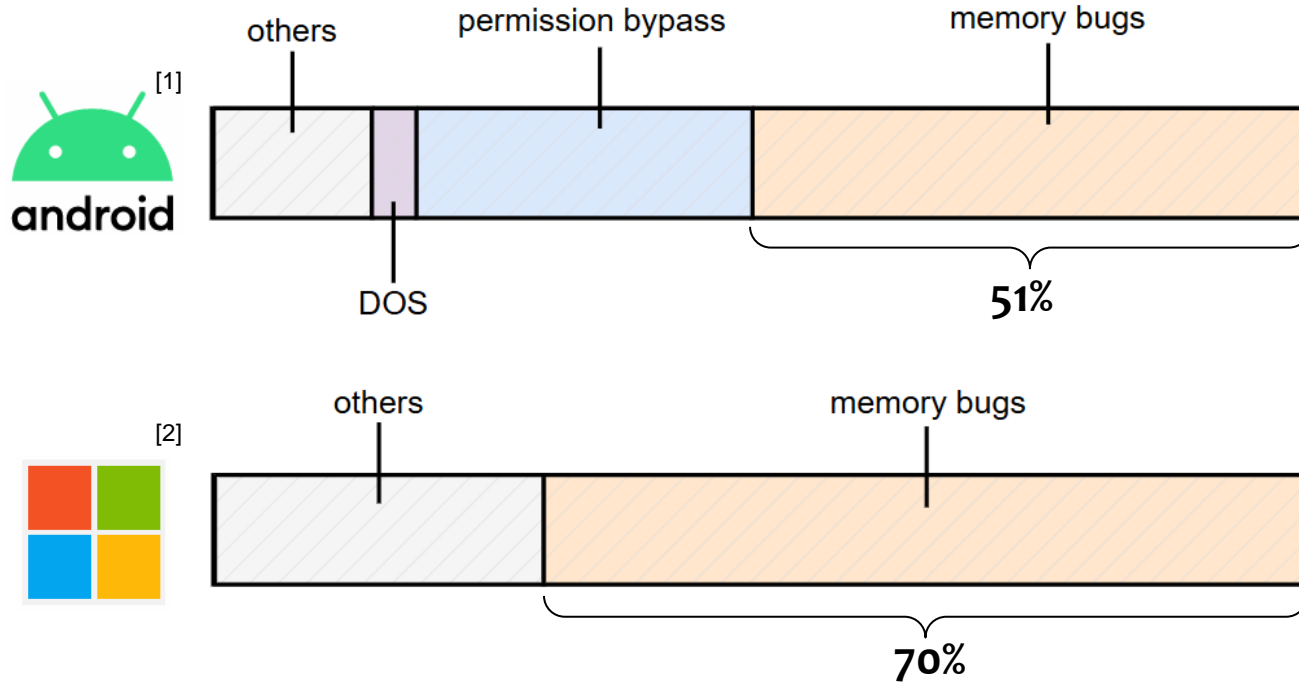
Outline



- Motivation
- Research Gap
- Problem Statement
- Background
- RECIPE for porting to CHERI
- Evaluation

- **Motivation**
- Research Gap
- Problem Statement
- Background
- RECIPE for porting to CHERI
- Evaluation

Memory Safety in Software, everywhere



[1] [Android Memory Safety](#)



[2] [Microsoft Memory Safety](#)

- **Spatial Memory Safety**

- **violation:** buffer overflow, stack overflow, etc.

- **Temporal Memory Safety**

- **violations:** use-after-free, double free, etc.



Outline



- Motivation
- **Research Gap**
- Problem Statement
- Background
- RECIPE for porting to CHERI
- Evaluation

- there are no memory safety solutions designed for DB systems:
 - Hashtables
 - Trees
- general software solutions:
 - ASLR, DEP, Stack Canaries, etc.
 - AddressSanitizer, Valgrind

- Motivation
- Research Gap
- **Problem Statement**
- Background
- RECIPE for porting to CHERI
- Evaluation

- **Databases** are susceptible to attacks that exploit memory vulnerabilities.
- **CHERI** tackles memory vulnerabilities by introducing fine-grained memory protection.
- the community did not focus on applying the CHERI to databases:
 - No RECIPE for porting complex data structures to CHERI
 - Unknown Overheads (Performance, Memory, etc.)

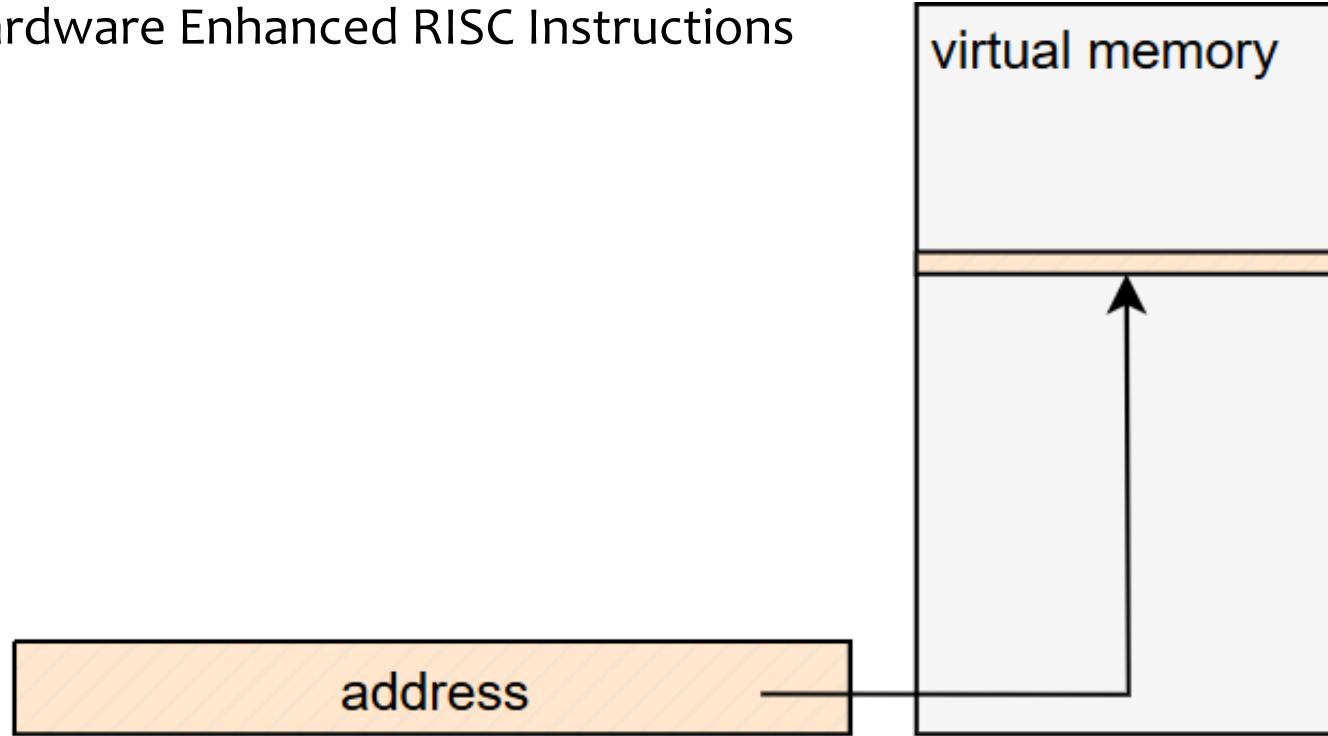
Outline



- Motivation
- Research Gap
- Problem Statement
- **Background**
- RECIPE for porting to CHERI
- Evaluation

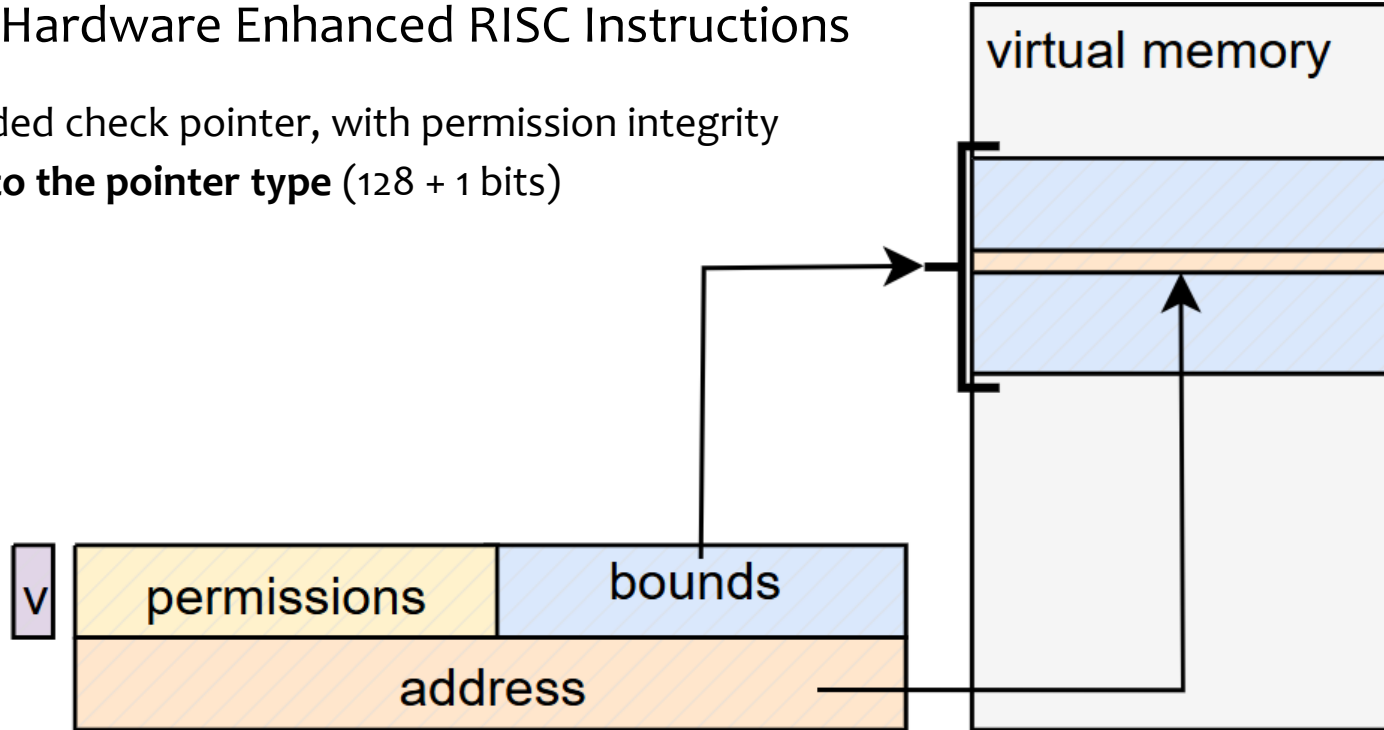
CHERI - Capability Hardware Enhanced RISC Instructions

CHERI - **Capability** Hardware Enhanced RISC Instructions



CHERI - **Capability** Hardware Enhanced RISC Instructions

- capability = bounded check pointer, with permission integrity
- **metadata added to the pointer type** (128 + 1 bits)



CHERI - Capability Hardware Enhanced RISC Instructions

- newly added instructions:
 - memory access (load & store) via boundary & permissions check
 - permissions management: decrease bounds, decrease permissions, etc.

CHERI faces disadvantages due to incompatibilities arising from the capability pointer format:

- pointers' capabilities of meta-data can interfere with pointer-arithmetic operations.
- pointers' size is extended, increasing the size of structures, breaking paddings, etc.

Outline



- Motivation
- Research Gap
- Problem Statement
- Background
- **RECIPE for porting to CHERI**
- Evaluation

- **Data structures** ported on CHERI:
 - Cache-Line Hash Table (CLHT): Lock-Based & Lock-Free versions
 - Cuckoo (libcuckoo library)
 - oneTBB (library) :: Concurrent Hash-Map (and subtypes)

 - Adaptive Radix Tree
 - Google BTree

RECIPE for porting to CHERI



1. Compiler & CHERI community
2. Base-Addresses used in Operations
 1. Pointer's Memory Address is used in Operations
 2. Primitive data is used interchangeably with pointer type
3. Source of provenance
4. Pointer Size Assumptions
 1. Updating the structures
 1. Re-Padding
 2. Re-Sizing
 2. Updating Functions

1. Compiler & CHERI community – compatibility with C++
 - standard libraries (boost, stdlib, etc.) are supported by the CHERI platform.
 - code that clearly separates pointer and primitive types does not require adjustments.

RECIPE for porting to CHERI

1. Compiler & CHERI community
2. Base-Addresses used in Operations
 - a. **Pointer's Memory Address is used in Operations**
 - used in: hash function computations, seed initialization, etc.

pointer



changes must not impact the hash, seed init.

```
get(void* address) {  
    return addresstable[hash(address)];  
}
```

RECIPE for porting to CHERI

1. Compiler & CHERI community
2. Base-Addresses used in Operations
 - a. **Pointer's Memory Address is used in Operations**
 - used in: hash function computations, seed initialization, etc.




```
get(void* address) {  
#ifdef __CHERI_PURE_CAPABILITY__  
    return addresstable[hash(cheri_getaddress(address))];  
#else  
    return addresstable[hash(address)];  
#endif  
}
```

RECIPE for porting to CHERI

1. Compiler & CHERI community
2. Base-Addresses used in Operations
 - a. Pointer's Memory Address is used in Operations
 - b. **Primitive data is used interchangeably with pointer type**
 - embedding information into pointer,


pointer



```
lock(void *address) {  
    state_type s = state.get();  
    if (!(s & BUSY)) {  
        state.update(address, WRITER);  
    }  
}
```

RECIPE for porting to CHERI

1. Compiler & CHERI community
2. Base-Addresses used in Operations
 - a. Pointer's Memory Address is used in Operations
 - b. **Primitive data is used interchangeably with pointer type**
 - embedding information into pointer,



```
lock(void *address) {  
    state_type s = state.get();  
#ifdef __CHERI_PURE_CAPABILITY__  
    if (!(cheri_getaddress(s) & BUSY)) {  
#else  
    if (!(s & BUSY)) {  
#endif  
        state.update(address, WRITER);  
    }  
}
```

RECIPE for porting to CHERI

1. Compiler & CHERI community
2. Base-Addresses used in Operations
3. Source of provenance

multiple pointers are involved in arithmetic operations.

```
void *foo() {  
    void *address = get_address();  
    void *mask     = get_coaddress();  
    return (address ^ mask);  
}
```

address pointer



mask pointer



result pointer



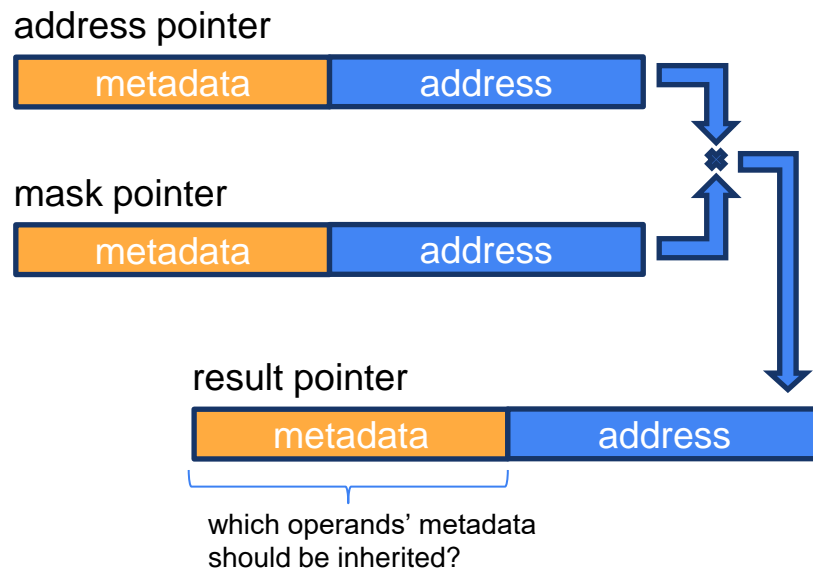
which operands' metadata
should be inherited?

RECIPE for porting to CHERI

1. Compiler & CHERI community
2. Base-Addresses used in Operations
3. Source of provenance



multiple pointers are involved in arithmetic operations.

```
void *foo() {  
    void *address = get_address();  
    void *mask    = get_coaddress();  
  
    /* breaking down the operation */  
    address ^= mask;  
    return address;  
}
```



RECIPE for porting to CHERI



1. Compiler & CHERI community
2. Base-Addresses used in Operations
3. Source of provenance
4. Pointer Size Assumptions
 - a. updating the structures
 - i. re-padding structures



RECIPE for porting to CHERI



1. Compiler & CHERI community
2. Base-Addresses used in Operations
3. Source of provenance
4. Pointer Size Assumptions
 - a. updating the structures
 - i. re-padding structures

aligned



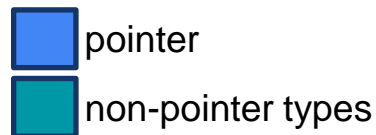

RECIPE for porting to CHERI

1. Compiler & CHERI community
2. Base-Addresses used in Operations
3. Source of provenance
4. Pointer Size Assumptions
 - a. updating the structures
 - i. re-padding structures



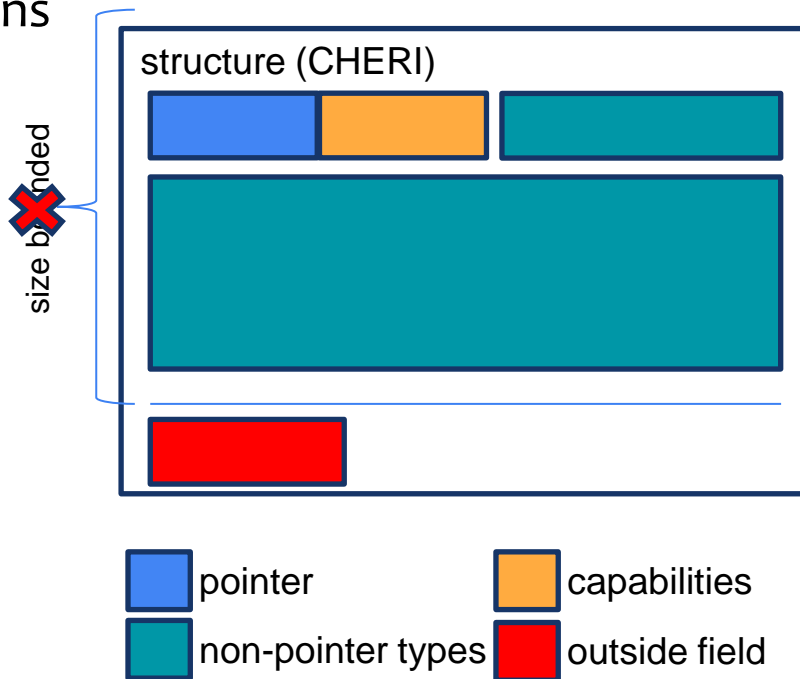
RECIPE for porting to CHERI

1. Compiler & CHERI community
2. Base-Addresses used in Operations
3. Source of provenance
4. Pointer Size Assumptions
 - a. **updating the structures**
 - i. re-padding structures
 - ii. **re-sizing structures**




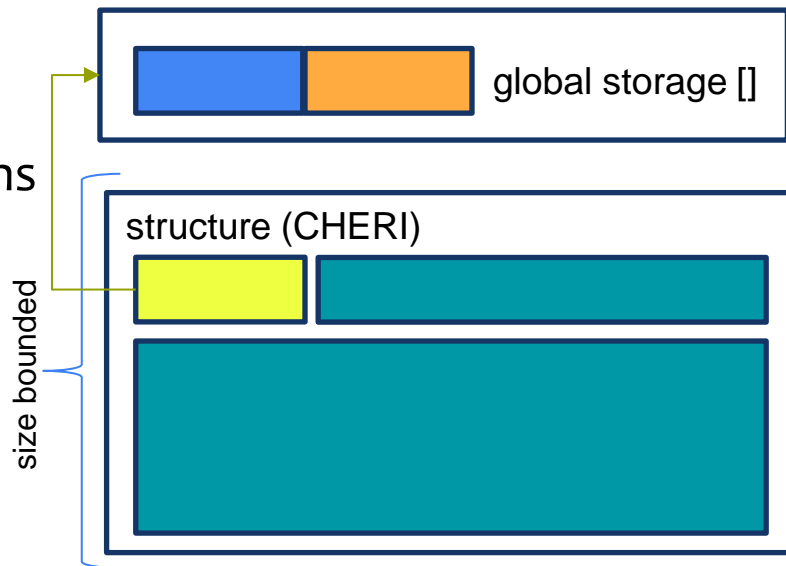
RECIPE for porting to CHERI

1. Compiler & CHERI community
2. Base-Addresses used in Operations
3. Source of provenance
4. Pointer Size Assumptions
 - a. **updating the structures**
 - i. re-padding structures
 - ii. **re-sizing structures**




RECIPE for porting to CHERI

1. Compiler & CHERI community
2. Base-Addresses used in Operations
3. Source of provenance
4. Pointer Size Assumptions
 - a. **updating the structures**
 - i. re-padding structures
 - ii. **re-sizing structures**



RECIPE for porting to CHERI

1. Compiler & CHERI community
2. Base-Addresses used in Operations
3. Source of provenance
4. Pointer Size Assumptions
 - a. updating the structures
 - i. re-padding structures
 - ii. re-sizing structures
 - b. **updating the functions, changing arguments**




```
void *foo(void* address) {  
    void *mem = allocate_mem(sizeof(address));  
    return mem;  
}
```


RECIPE for porting to CHERI

1. Compiler & CHERI community
2. Base-Addresses used in Operations
3. Source of provenance
4. Pointer Size Assumptions
 - a. updating the structures
 - i. re-padding structures
 - ii. re-sizing structures
 - b. updating the functions, changing arguments

pointer



```
#ifdef __CHERI_PURE_CAPABILITY__  
void *foo(uint64_t address) {  
#else  
void *foo(void* address) {  
#endif  
    void *mem = allocate_mem(sizeof(address));  
    return mem;  
}
```





- **Custom Framework** was developed for evaluating the data-structures
 - designed microbenchmarks to evaluate the latency, memory and PMCstats for all data-structure operations – insert, query and delete.
 - used a modified version of YCSB to support data-structure evaluation.
- **Evaluation**
 - performed using ARM (without capabilities) as baseline.
 - experiments performed on 100.000 operations.
 - filling factor of 95%.

1. Latency – Microbenchmarks, Macrobenchmarks (YCSB)
 1. How does CHERI impact the execution time of DBMS operations?
2. Memory Overhead – Microbenchmarks
 1. How much extra memory is required to store capabilities compared to traditional pointers?
3. Cache Usage – Microbenchmarks
 1. Does the larger size of capabilities affect cache efficiency?

Evaluation :: Latency Overhead

Insertion Operation:





- hashtables: 40% - 90%
- trees: 95% - 130%



Evaluation :: Latency Overhead

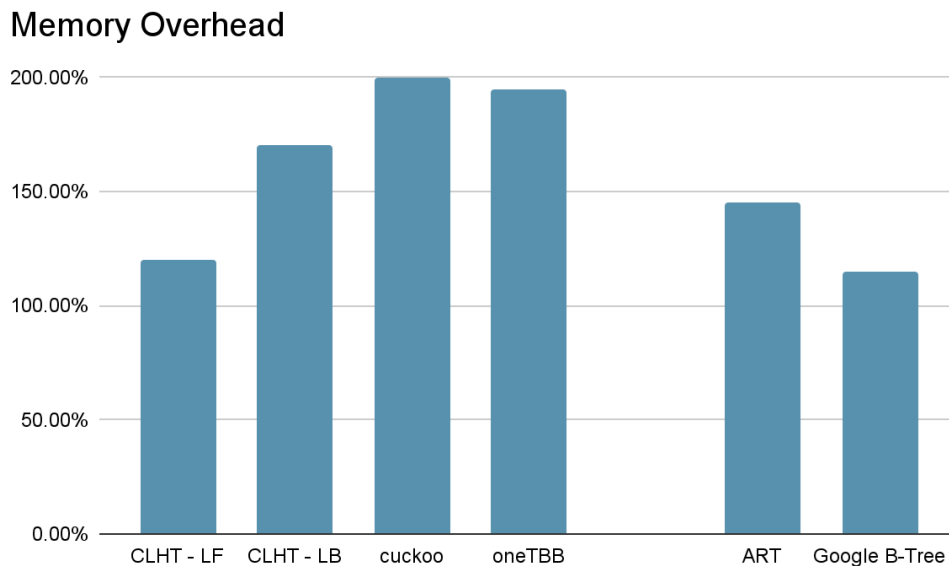
Query Operation:

- hashtables: 70% - 80%
- trees: 50% - 120%



Evaluation :: Memory Overhead

- hashtables: 120% - 200%
- trees: 115% - 150%




Data Structure	Cache L1 (Miss / Read) Overhead
Cuckoo	~ 1.0x
oneTBB Concurrent Hashtable	~ 1.2-2.0x
CLHT Lock-Free	~ 3.0-5.0x
CLHT Lock-Based	~ 3.0-5.0x
Adaptive Radix Tree	~ 1.1-1.2x
Google BTree	~ 1.1-1.2x

- Databases are vulnerable to memory-based exploits.
 - CHERI offers a solution to mitigate the memory-safe problem.
- CHERI requires careful consideration during the porting procedure.
 - CHERI uses larger pointers to accommodate metadata for memory protection.
 - it impacts data structure, existing code, and external-library usage.
 - existing code often needs adjustments to handle these larger pointers.
- CHERI technology introduces performance overheads.
 - CHERI's security benefits come with potential performance trade-offs in: memory management, latency, and cache efficiency.

Backup

Evaluation :: Memory Overhead

- hashtables: 120% - 200%
- trees: 115% - 150%



Evaluation :: Performance Overhead

Insertion Operation:

- hashtables: 35% - 140%
- trees: 95% - 130%

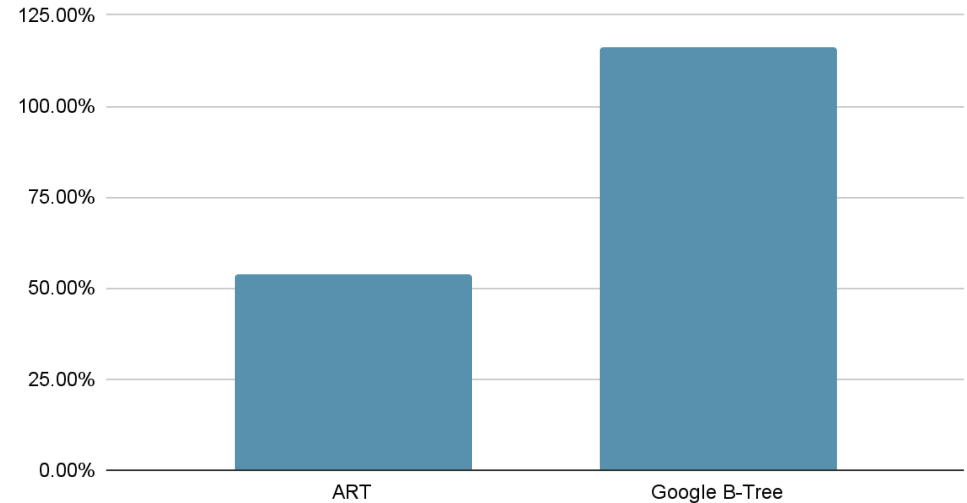
Query Operation:

- hashtables: 70% - 80%
- trees: 50% - 120%

Deletion Operations:

- trees: 55% - 125%

Performance Overhead Deletion



Evaluation :: Cache Overhead

Data Structure	PMC Stats			
	Stat Name	No Capability	Capability	Increase
Cuckoo Hashtable	LL read	29.2 M.	27.5 M.	0.941x
	LL miss	18.0 M.	15.9 M.	0.883x
	L1 read	4772.8 M.	4726.4 M.	0.99x
	L1 write	2801.8 M.	3522.2 M.	1.257x
OneTBB Conc. HT.	LL read	2.5 M.	2.9 M.	1.16x
	LL miss	2.2 M.	2.7 M.	1.227x
	L1 read	1160.3 M.	1284.8 M.	1.107x
	L1 write	868.0 M.	1002.7 M.	1.155x
CLHT Lock-Base	LL read	86.7 M.	321.9 M.	3.712x
	LL miss	54.2 M.	248.7 M.	4.588x
	L1 read	1861.2 M.	1810.7 M.	0.972x
	L1 write	841.1 M.	854.0 M.	1.015x
CLHT Lock-Free	LL read	1.8 M.	3.1 M.	1.722x
	LL miss	1.7 M.	3.1 M.	1.823x
	L1 read	749.1 M.	1177.4 M.	1.571x
	L1 write	542.0 M.	902.2 M.	1.664x

Evaluation :: Cache Overhead

Data Structure	PMC Stats			
	Stat Name	No Capability	Capability	Increase
Adaptive Radix Tree	LL read	3.0 M.	3.8 M.	1.266x
	LL miss	2.9 M.	3.8 M.	1.31x
	L1 read	1497.5 M.	1651.9 M.	1.103x
	L1 write	1056.7 M.	1222.1 M.	1.156x
Google BTree	LL read	3.0 M.	3.6 M.	1.2x
	LL miss	3.0 M.	3.6 M.	1.2x
	L1 read	1725.6 M.	1884.8 M.	1.092x
	L1 write	1238.1 M.	1416.3 M.	1.143x