


Serverless computing in unikernels using eBPF code injection

Author: Milen Vitanov
Supervisor: Prof. Dr. Pramod Bhatotia
Advisors: Ilya Meignan-Masson
Masanori Misono




TUM Uhrenturm

01.06.2024 - 01.12.2024

Outline

- 1 Background & Motivation
- 2 Design
- 3 Implementation Details
- 4 Evaluation
- 5 Conclusion & Future Work

Virtualization Technologies

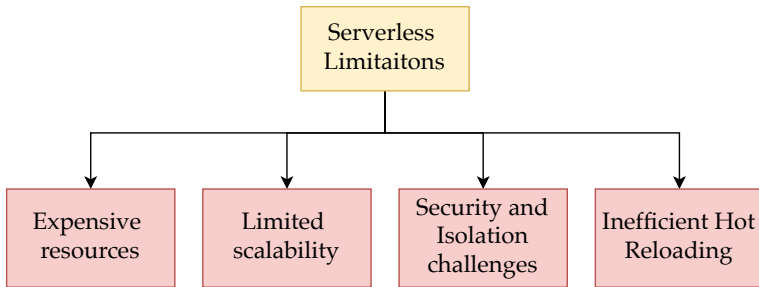


Serverless

- Event-Driven Execution
 - No Server Management, Pay-Per-Use Pricing
 - Automatic Scaling
-
- Data Processing: **MapReduce**

eBPF in Unikernels

- Dynamic Functionality
- Performance Monitoring
- Enhanced Security



■ General-purpose OS

■ Context Switches
■ Initialization Time

■ Lack of isolation
■ Shared resources

■ Secure code reloading

Motivation Problem

How can we build a performant serverless data processing framework which allocates only resources that it needs?

Existing issues

- slow initialization time
- a lot of resources allocated for the OS
- limited isolation of distributed workers


Motivation Overview

Proposal

Design a new serverless data processing method using unikernels as workers, equipped with eBPF dynamic code injection to allow for any serverless use case.

Our framework should:

- use resources efficiently
- be performant
- offer high security & isolation
- be able to inject code at runtime
- be flexible and usable.

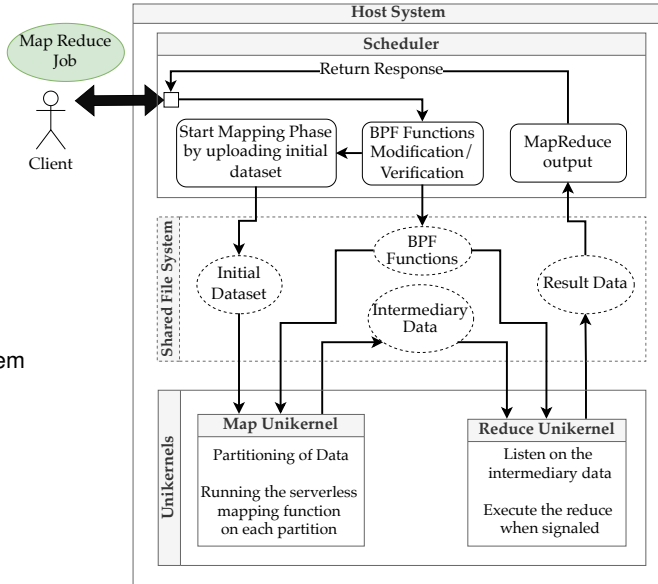


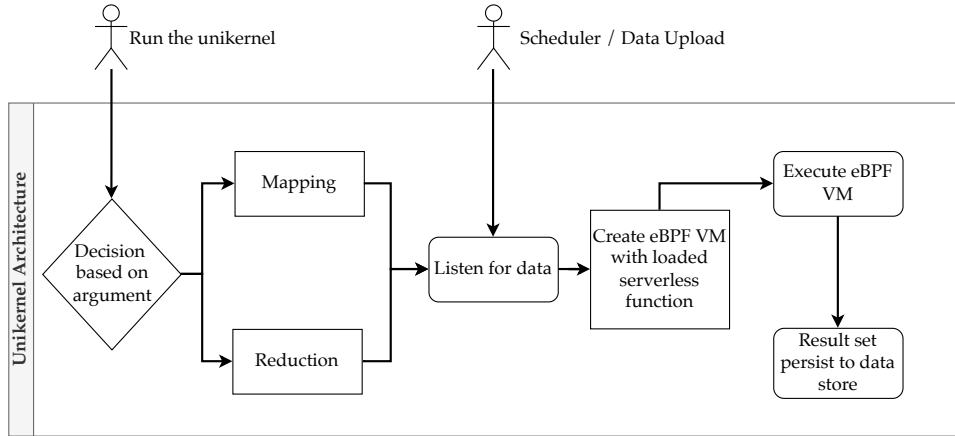
Outline

- 1 Background & Motivation
- 2 Design**
- 3 Implementation Details
- 4 Evaluation
- 5 Conclusion & Future Work

Design Components

- Scheduler
- Shared File System
- Unikernels





■ **Challenge I:** Integrating unikernels and eBPF into the data processing algorithm

- ☐ Work with concrete addresses, simple and small stack memory usage
- ☐ Predefined BPF hooks according to the execution model
- ☐ File System Synchronization Level

■ **Challenge II:** Defining the BPF Helper Functions

- ☐ Basic read/write operations
- ☐ Generic communication data structure
- ☐ Generic Helpers API

■ **Challenge III:** Complex BPF Verification

- ☐ Decoupled Verification using a robust external verifier

Challenge I

Integrating unikernels and eBPF into the data processing algorithm

- uBPF Integration - porting of an existing library
- Fixed-role structure of the unikernel workers
 - More structured and predictable system
- Mapper Phase Hook
- Synchronization of workers interaction
 - direct and efficient data flow
 - low latency
- Reducer Phase Hook

Challenge II

Defining the BPF Helper Functions

BPF Limitations

- No pointer arithmetics
- Limited stack size
- No floating point numbers
- No dynamic data structures
- No unbounded loops
- Limited/Non-existent Debugging

Data Structure Helpers

- Read/Modify the transformed data
- Traverse the data
- Allocate/deallocate the structure

Data Helpers

- Read/Modify the incoming dataset
- Traverse the dataset

Debugging Helpers

- Allow for printing messages/numbers
- Allow for printing the transformed data

Challenge II

Data Structure Helpers API \Rightarrow `bpf_hashtable_...`

#	Helper Name	Description
1	<code>bpf_hashtable_initialize</code>	Initialize new hashtable in host
2	<code>bpf_hashtable_free</code>	Deallocate hashtable
3	<code>bpf_hashtable_get_value_address</code>	Retrieve value address linked to a key
4	<code>bpf_hashtable_get_map_key_value_address</code>	Retrieve value address linked to a key, which is a value in another map
5	<code>bpf_hashtable_get_data_key_value_address</code>	Retrieve value address linked to a key, which is a substring from data at offset
6	<code>bpf_hashtable_insert_value</code>	Insert a value for a key
7	<code>bpf_hashtable_insert_map_key_value</code>	Insert a value for a key, which is a value to another map key
8	<code>bpf_hashtable_insert_data_key_value</code>	Insert a value for a key, which is a substring from data at offset
9	<code>bpf_hashtable_get_size</code>	Retrieve hashtable bucket size
10	<code>bpf_hashtable_get_first_node</code>	Retrieve first node address at bucket index
11	<code>bpf_hashtable_get_next_node</code>	Retrieve next linked node to a given node
12	<code>bpf_hashtable_get_node_value_size</code>	Retrieve the value size of a node
13	<code>bpf_hashtable_get_node_value</code>	Retrieve the value of a node
14	<code>bpf_hashtable_get_node_key</code>	Retrieve the key of a node
15	<code>bpf_hashtable_insert_hashtable</code>	Inserts a map as a value inside another map; key is a string
16	<code>bpf_hashtable_insert_map_key_hashtable</code>	Inserts a map as a value in another map; key is linked val to key is another map
17	<code>bpf_hashtable_insert_data_key_hashtable</code>	Inserts a map as a value in another map; key linked to a key, which is a substring from data at offset
...27	<code>..._node_address, ..._int</code>	simplified functions when working directly with integers or nodes

Challenge II

Data Helpers API \Rightarrow `bpf_data_...`

#	Helper Name	Description
1	<code>bpf_data_get_next_index_delimiter</code>	Retrieve the next index pointing to a matching delimiter in the data chunk
2	<code>bpf_data_get_next_index_map_word</code>	Retrieve the next index pointing to a word linked to a map key
3	<code>bpf_data_update_index_content</code>	Update the data at index with a new byte content

Challenge II

Debugging Helpers API \Rightarrow `bpf_print_...`

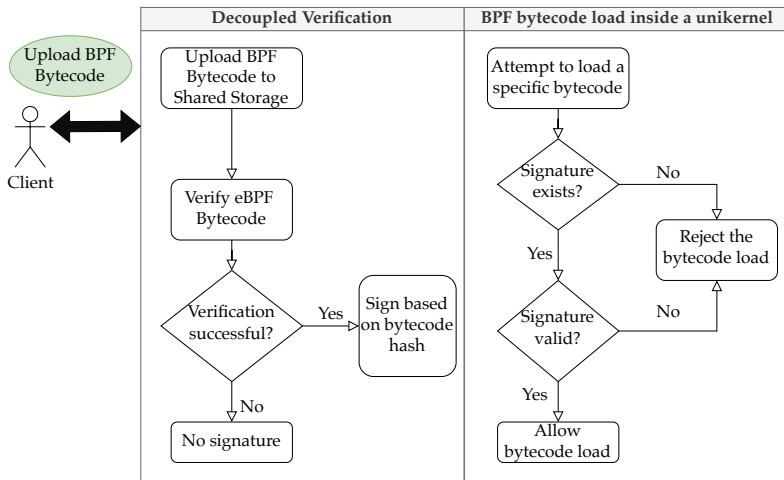
#	Helper Name	Description
1	<code>bpf_print_text</code>	Print a message up to a specified size
2	<code>bpf_print_integer</code>	Prints a 64-bit integer
3	<code>bpf_print_hashtable</code>	Prints a table content

Decoupled Verification [2]

- Reduced complexity
 - Unikernel is kept minimal and optimized for its task
- Flexibility
 - More powerful system can run the verification for resource-intensive analysis
- Scalability
 - In multi-node environment, single centralized verifier can service multiple unikernels reducing duplication

Challenge III

Verification Design



Outline

- 1 Background & Motivation
- 2 Design
- 3 Implementation Details**
- 4 Evaluation
- 5 Conclusion & Future Work

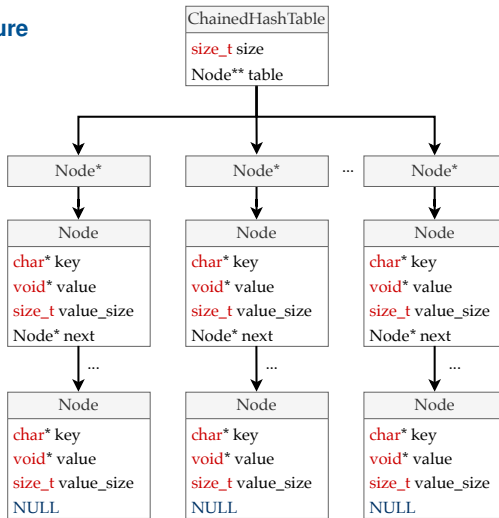
Implementation Details

External projects

- Unikernels: Unikraft [7]
 - easy to configure, optimized initialization time; modular library approach
- BPF library: uBPF [9]
 - there is already a porting unikraft library done by V. Hendrychová [6]
- BPF Verification: PREVAIL [10]
 - strict, robust and already used in operating systems like Windows
- Data processing algorithm: MapReduce [3]
 - fundamental data processing execution model at the core of the majority of industry-wide solutions
- Virtual Emulator: QEMU [4]
 - popular, with extensive documentation and appropriate to simulate different VM scenarios
- Shared File System: 9pfs [5]
 - lightweight, suitable for minimalist environments
- Cryptography: OpenSSL [8]

Implementation Details

Communication Data Structure



Implementation Details

BPF Program Context

struct Context
<code>uint64_t data_address</code>
<code>uint64_t data_size</code>
<code>uint64_t data_structure_address</code>

Pointer cast to uint64_t

```
char* line_ptr = (char*)malloc(DYNAMIC_SIZE * sizeof(char));  
context->data_address = (uint64_t)(uintptr_t) line_ptr;
```

uint64_t cast to a pointer

```
char* line_ptr = (char*)(uintptr_t) context->data_address;
```

Implementation Details

PREVAIL Custom Platform

eBPF Helper Prototype in PREVAIL

```
static const struct EbpfHelperPrototype bpf_print_text_proto = {  
    .name = "print_text",  
    .return_type = EBPF_RETURN_TYPE_INTEGER_OR_NO_RETURN_IF_SUCCEED,  
    .argument_type = {  
        EBPF_ARGUMENT_TYPE_PTR_TO_READABLE_MEM,  
        EBPF_ARGUMENT_TYPE_CONST_SIZE  
    },  
};
```

PREVAIL Context Descriptor

```
constexpr ebpf_context_descriptor_t bpf_prog_descr = {  
    .size = 256,  
    .data = 0,  
    .end = 255,  
    .meta = -1  
};
```

Outline

- 1 Background & Motivation
- 2 Design
- 3 Implementation Details
- 4 Evaluation**
- 5 Conclusion & Future Work

Qualitative Analysis



- Development Complexity
- Debugging
- Learning Curve
- Code Maintenance

Quantitative Evaluation

- Initialization Time
- Performance
- Memory Footprint
- CPU-Intensive BPF Functions
- Flexibility




Evaluation

Initialization Time






Unique Word Count

Evaluation Performance



Evaluation

Memory Footprint




Evaluation

CPU-Intensive BPF Functions

- Mapper: Fibonacci number calculation ($N = 50000$)
- Reducer: π calculation using the Leibniz Series ($I = 1000000000$, $S = 100000000$)

	Average across 10 runs
Mapper Unikernel	133.877MB
Reducer Unikernel	133.855MB

Memory Footprint




Evaluation

Flexibility: Search Engine

BPF Mapper Function

```
for (i : data_size) {  
    int token_start = i;  
    int match_idx = bpf_data_get_next_index_map_word(data,  
        data_size, ds_address, MAP_SEARCH_KEY, i);  
  
    if (match_idx = data_size) { // no match  
        break;  
    }  
  
    int value = bpf_hashtable_get_map_key_int(  
        ds_address, MAP_DOC_ID_KEY);  
  
    bpf_hashtable_insert_map_key_int(  
        ds_address, MAP_DOC_ID_KEY, ++value);  
  
    i = match_idx;  
}
```



nonBPF again slightly slower

Transformed Data

$word_1 \rightarrow [(doc_id_1, occurrences), \dots (doc_id_n, occurrences)]$

$word_2 \rightarrow [(doc_id_2, occurrences), \dots (doc_id_k, occurrences)]$

...

$word_m \rightarrow [...]$

Outline

- 1 Background & Motivation
- 2 Design
- 3 Implementation Details
- 4 Evaluation
- 5 Conclusion & Future Work**

Serverless data processing framework built on unikernels with the ability to dynamically inject code

- Better initialization time
- Faster performance with bigger datasets
- Efficient memory footprint
- Capable of CPU-intensive tasks
- Lagging behind Linux due to missing unikernel internal optimizations
- Flexible Helpers API

Bibliography I

- [1] Mihnea Firoiu Alexander Jung. *Unikernel Concepts*. URL: <https://unikraft.org/docs/concepts>.
- [2] Milo Craun, Adam Oswald, and Dan Williams. “Enabling eBPF on Embedded Systems Through Decoupled Verification”. In: *Proceedings of the 1st Workshop on EBPF and Kernel Extensions*. eBPF '23. New York, NY, USA: Association for Computing Machinery, 2023, 63–69. ISBN: 9798400702938. DOI: 10.1145/3609021.3609299. URL: <https://doi.org/10.1145/3609021.3609299>.
- [3] Jeffrey Dean and Sanjay Ghemawat. “MapReduce: simplified data processing on large clusters”. In: *Commun. ACM* 51.1 (2008), 107–113. ISSN: 0001-0782. DOI: 10.1145/1327452.1327492. URL: <https://doi.org/10.1145/1327452.1327492>.
- [4] The QEMU Project Developers. *QEMU Docs*. URL: <https://www.qemu.org/docs/master/>.
- [5] Plan 9 Foundation. *Plan 9 from Bell Labs*. URL: <https://9p.io/plan9/>.
- [6] Vanda Hendrychová. *ubpf for Unikraft*. Accessed: 2024-06-15. 2023. URL: https://github.com/vandah/uk_ubpf.

Bibliography II

- [7] Simon Kuenzer et al. “Unikraft: fast, specialized unikernels the easy way”. In: *Proceedings of the Sixteenth European Conference on Computer Systems*. EuroSys '21. Online Event, United Kingdom: Association for Computing Machinery, 2021, 376–394. ISBN: 9781450383349. DOI: 10.1145/3447786.3456248. URL: <https://doi.org/10.1145/3447786.3456248>.
- [8] 2024 OpenSSL. *OpenSSL*. URL: <https://www.openssl.org/>.
- [9] Iovisor Team. *uBPF: Userspace eBPF VM*. Accessed: 2024-10-15. 2024. URL: <https://github.com/charlespwd/project-title>.
- [10] PREVAIL Team. *PREVAIL - A new eBPF verifier*. Accessed: 2024-11-15. 2024. URL: <https://github.com/vbpf/ebpf-verifier>.